Big Data Processing Lab-5

Name: Sumit Vishwakarma

Roll No: 202201320

Colab Link:https://colab.research.google.com/drive/1qoWgXo2rUcqI-ZVKYnXp_uRl_PazjuW8 - scrollTo=YKAMqYiUM7cK

Q1: Add a column ID in both the files that identify the class in the class file and the data record in the data vector file.

Code and Output:

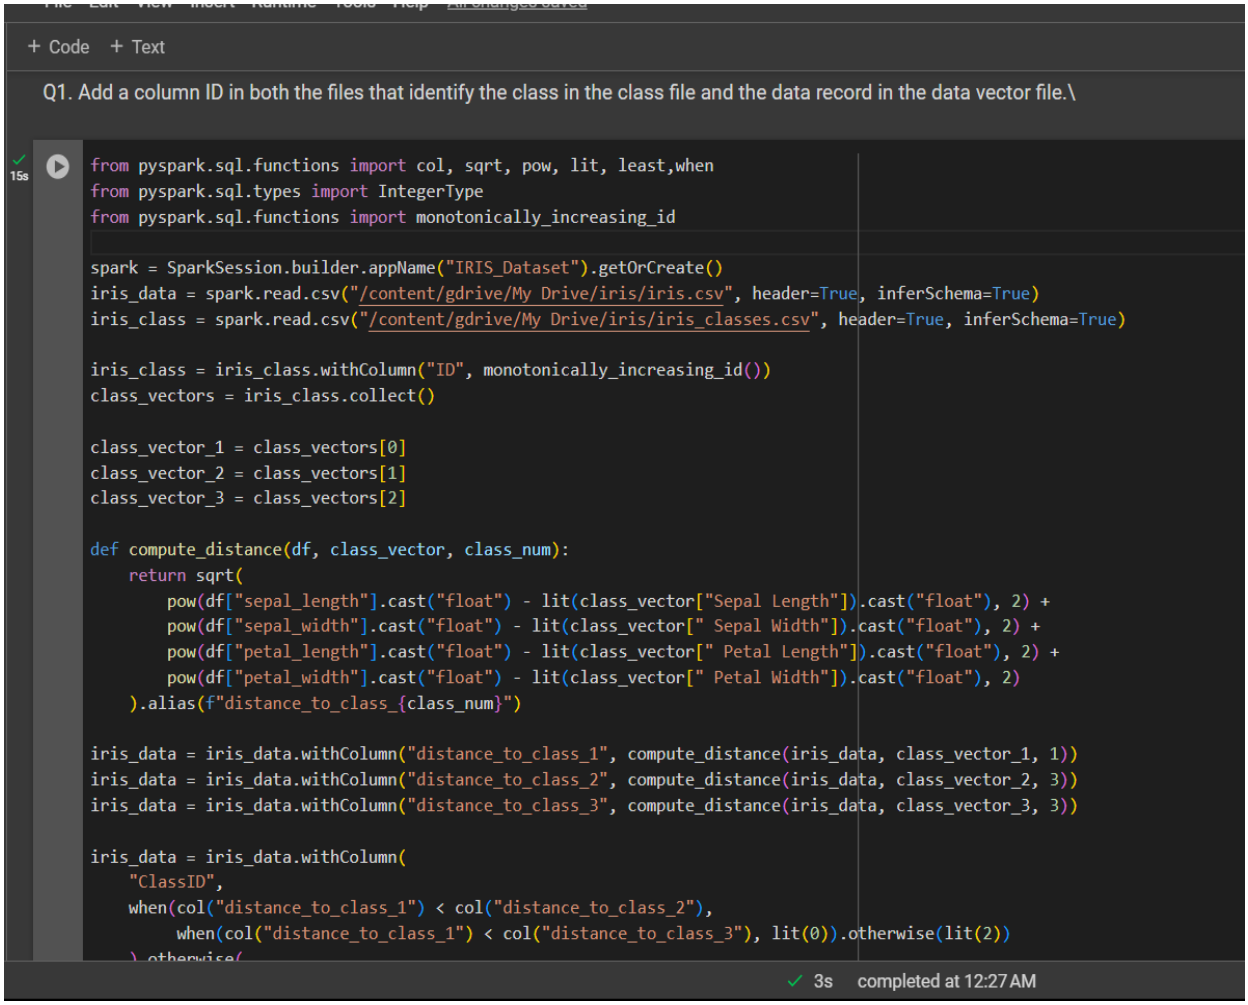

+ Code + Text

Q1. Add a column ID in both the files that identify the class in the class file and the data record in the data vector file.\

```
from pyspark.sql.functions import col, sqrt, pow, lit, least,when
from pyspark.sql.types import IntegerType
from pyspark.sql.functions import monotonically_increasing_id

spark = SparkSession.builder.appName("IRIS_Dataset").getOrCreate()
iris_data = spark.read.csv("/content/gdrive/My Drive/iris/iris.csv", header=True, inferSchema=True)
iris_class = spark.read.csv("/content/gdrive/My Drive/iris/iris_classes.csv", header=True, inferSchema=True)

iris_class = iris_class.withColumn("ID", monotonically_increasing_id())
class_vectors = iris_class.collect()

class_vector_1 = class_vectors[0]
class_vector_2 = class_vectors[1]
class_vector_3 = class_vectors[2]

def compute_distance(df, class_vector, class_num):
    return sqrt(
        pow(df["sepal_length"].cast("float") - lit(class_vector["Sepal Length"]).cast("float"), 2) +
        pow(df["sepal_width"].cast("float") - lit(class_vector[" Sepal Width"]).cast("float"), 2) +
        pow(df["petal_length"].cast("float") - lit(class_vector[" Petal Length"]).cast("float"), 2) +
        pow(df["petal_width"].cast("float") - lit(class_vector[" Petal Width"]).cast("float"), 2)
    ).alias(f"distance_to_class_{class_num}")

iris_data = iris_data.withColumn("distance_to_class_1", compute_distance(iris_data, class_vector_1, 1))
iris_data = iris_data.withColumn("distance_to_class_2", compute_distance(iris_data, class_vector_2, 3))
iris_data = iris_data.withColumn("distance_to_class_3", compute_distance(iris_data, class_vector_3, 3))

iris_data = iris_data.withColumn(
    "ClassID",
    when(col("distance_to_class_1") < col("distance_to_class_2"),
         when(col("distance_to_class_1") < col("distance_to_class_3"), lit(0)).otherwise(lit(2))
```

15s

✓ 3s completed at 12:27 AM

```
+------------+-----------+------------+-----------+-------+
|sepal_length|sepal_width|petal_length|petal_width|ClassID|
+------------+-----------+------------+-----------+-------+
|         5.1|        3.5|         1.4|        0.2|      1|
|         4.9|        3.0|         1.4|        0.2|      1|
|         4.7|        3.2|         1.3|        0.2|      1|
|         4.6|        3.1|         1.5|        0.2|      1|
|         5.0|        3.6|         1.4|        0.2|      1|
|         5.4|        3.9|         1.7|        0.4|      1|
|         4.6|        3.4|         1.4|        0.3|      1|
|         5.0|        3.4|         1.5|        0.2|      1|
|         4.4|        2.9|         1.4|        0.2|      1|
|         4.9|        3.1|         1.5|        0.1|      1|
|         5.4|        3.7|         1.5|        0.2|      1|
|         4.8|        3.4|         1.6|        0.2|      1|
|         4.8|        3.0|         1.4|        0.1|      1|
|         4.3|        3.0|         1.1|        0.1|      1|
|         5.8|        4.0|         1.2|        0.2|      1|
|         5.7|        4.4|         1.5|        0.4|      1|
|         5.4|        3.9|         1.3|        0.4|      1|
|         5.1|        3.5|         1.4|        0.3|      1|
|         5.7|        3.8|         1.7|        0.3|      1|
|         5.1|        3.8|         1.5|        0.3|      1|
+------------+-----------+------------+-----------+-------+
only showing top 20 rows

+------------+------------+-------------+------------+---+
|Sepal Length| Sepal Width| Petal Length| Petal Width| ID|
+------------+------------+-------------+------------+---+
|      5.8885|      2.7377|       4.3967|       1.418|  0|
|       5.006|       3.418|        1.464|       0.244|  1|
|      6.8462|      3.0821|       5.7026|      2.0795|  2|
+------------+------------+-------------+------------+---+
```

✓ 3s completed at 12:27 AM

Q2: Create two tables iris_data and iris_class from respective files.

Code and Output:

+ Code + Text

Q2. Create two tables iris_data and iris_class from respective files.

```
iris_data.createOrReplaceTempView("iris_data")
iris_class.createOrReplaceTempView("iris_class")
spark.sql("SELECT * FROM iris_data").show()
spark.sql("SELECT * FROM iris_class").show()
```

```
+------------+-----------+------------+-----------+-------------------+--------------------+--------------------+-------+
|sepal_length|sepal_width|petal_length|petal_width|distance_to_class_1| distance_to_class_2|distance_to_class_3|ClassID|
+------------+-----------+------------+-----------+-------------------+--------------------+--------------------+-------+
|         5.1|        3.5|         1.4|        0.2|  3.415635759683184| 0.14694209172047115|  5.026802273454714|      1|
|         4.9|        3.0|         1.4|        0.2| 3.3925900687887354| 0.43816888941238546|  5.083239154556709|      1|
|         4.7|        3.2|         1.3|        0.2| 3.5636117914381154|  0.4123010130666758|   5.24742650008455|      1|
|         4.6|        3.1|         1.5|        0.2| 3.4155363358648554|  0.5188373176513053|  5.122509339134511|      1|
|         5.0|        3.6|         1.4|        0.2| 3.4636292199923977| 0.19796962736760818| 5.0716427988887665|      1|
|         5.4|        3.9|         1.7|        0.4| 3.1461227005790016|  0.6838071222054748| 4.6477929732942895|      1|
|         4.6|        3.4|         1.4|        0.3| 3.5112687595017684|  0.4152012883299955|  5.179330225193297|      1|
|         5.0|        3.4|         1.5|        0.2|  3.332036651820808|0.059933274503142425|   4.9702999532176|      1|
|         4.4|        2.9|         1.4|        0.2|   3.56450663081957|  0.8009942696481962|  5.297352182415964|      1|
|         4.9|        3.1|         1.5|        0.1| 3.3520721873213133|  0.36659519580872313|  5.036693480164783|      1|
|         5.4|        3.7|         1.5|        0.2|  3.322506276766178| 0.48784431901928277|  4.864942076450807|      1|
|         4.8|        3.4|         1.6|        0.2| 3.3058324929805014|  0.2513800802567068|  4.965062096506971|      1|
|         4.8|        3.0|         1.4|        0.1| 3.4599085595947954|  0.4919267464898315|  5.159889530578086|      1|
|         4.3|        3.0|         1.1|        0.1| 3.8983956196432366|  0.9090609147463025|  5.620695816912593|      1|
|         5.8|        4.0|         1.2|        0.2|  3.647416538177095|   1.020192248542253|  5.073755878027691|      1|
|         5.7|        4.4|         1.5|        0.4| 3.4965652307950275|  1.2130918761572562| 4.8511032471111175|      1|
|         5.4|        3.9|         1.3|        0.4|  3.495060589448569|  0.6624138908804096|  4.996405141091827|      1|
|         5.1|        3.5|         1.4|        0.3| 3.3812671987049567| 0.15097012740350055|  4.990274687166005|      1|
|         5.7|        3.8|         1.7|        0.3|  3.112254458226747|  0.8284875880413031|  4.584384373299028|      1|
|         5.1|        3.8|         1.5|        0.3| 3.3750567946218015|  0.3989886625658477| 4.9389333115940195|      1|
+------------+-----------+------------+-----------+-------------------+--------------------+--------------------+-------+
```

✓ 3s completed at 12:27 AM

Q3: Implement a user-defined function called classify_iris where a data vector is input and returns an integer representing a class number (1 to 3).

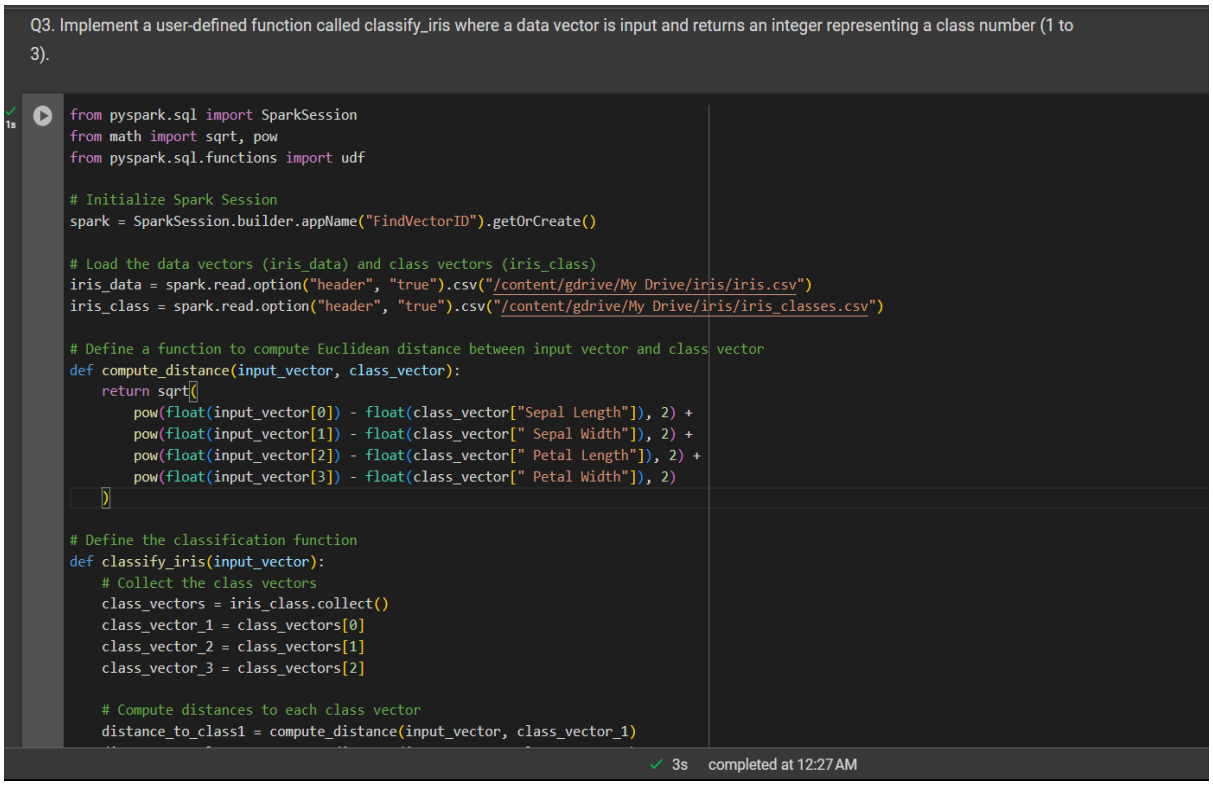

Q3. Implement a user-defined function called classify_iris where a data vector is input and returns an integer representing a class number (1 to 3).

```
from pyspark.sql import SparkSession
from math import sqrt, pow
from pyspark.sql.functions import udf

# Initialize Spark Session
spark = SparkSession.builder.appName("FindVectorID").getOrCreate()

# Load the data vectors (iris_data) and class vectors (iris_class)
iris_data = spark.read.option("header", "true").csv("/content/gdrive/My Drive/iris/iris.csv")
iris_class = spark.read.option("header", "true").csv("/content/gdrive/My Drive/iris/iris_classes.csv")

# Define a function to compute Euclidean distance between input vector and class vector
def compute_distance(input_vector, class_vector):
    return sqrt(
        pow(float(input_vector[0]) - float(class_vector["Sepal Length"]), 2) +
        pow(float(input_vector[1]) - float(class_vector[" Sepal Width"]), 2) +
        pow(float(input_vector[2]) - float(class_vector[" Petal Length"]), 2) +
        pow(float(input_vector[3]) - float(class_vector[" Petal Width"]), 2)
    )

# Define the classification function
def classify_iris(input_vector):
    # Collect the class vectors
    class_vectors = iris_class.collect()
    class_vector_1 = class_vectors[0]
    class_vector_2 = class_vectors[1]
    class_vector_3 = class_vectors[2]

    # Compute distances to each class vector
    distance_to_class1 = compute_distance(input_vector, class_vector_1)
```

✓ 3s completed at 12:27 AM

Q4: Use this user function to classify all data vectors using an SQL statement SELECT ID, classify_iris (*) from iris_data

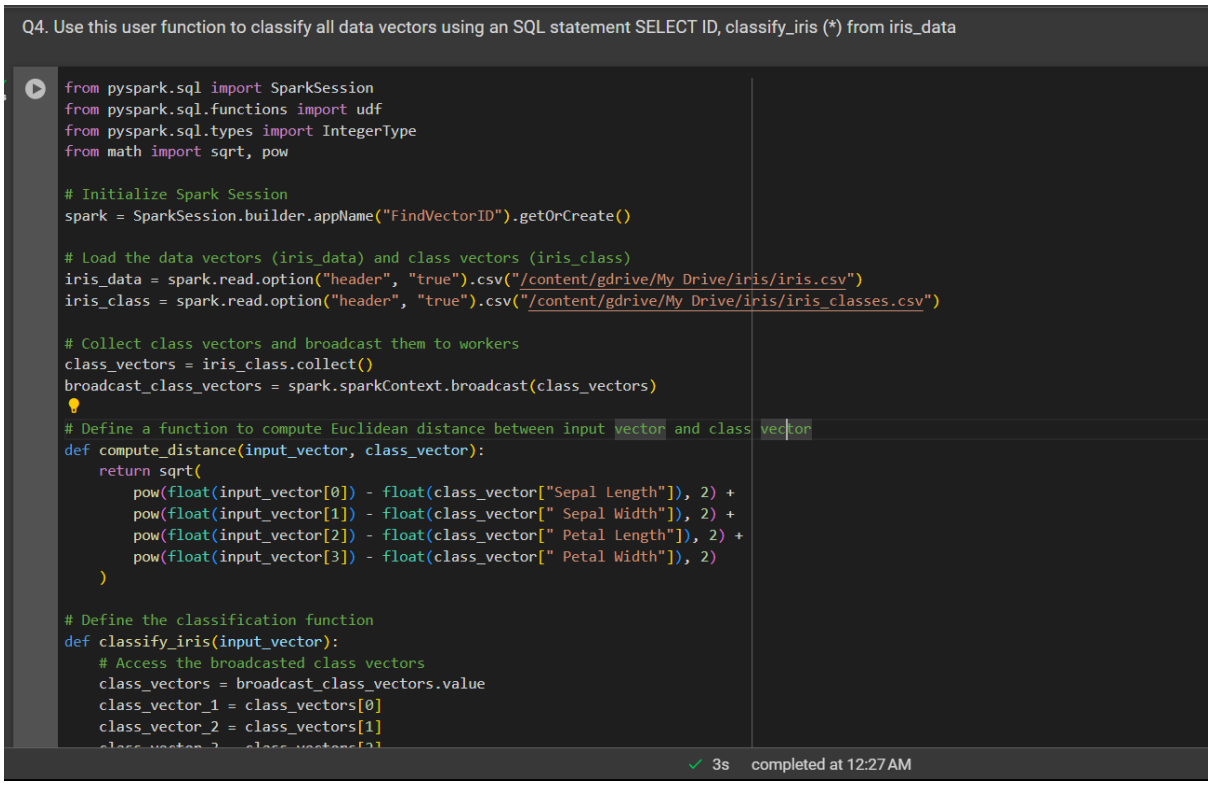

Q4. Use this user function to classify all data vectors using an SQL statement SELECT ID, classify_iris (*) from iris_data

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import udf
from pyspark.sql.types import IntegerType
from math import sqrt, pow

# Initialize Spark Session
spark = SparkSession.builder.appName("FindVectorID").getOrCreate()

# Load the data vectors (iris_data) and class vectors (iris_class)
iris_data = spark.read.option("header", "true").csv("/content/gdrive/My Drive/iris/iris.csv")
iris_class = spark.read.option("header", "true").csv("/content/gdrive/My Drive/iris/iris_classes.csv")

# Collect class vectors and broadcast them to workers
class_vectors = iris_class.collect()
broadcast_class_vectors = spark.sparkContext.broadcast(class_vectors)

# Define a function to compute Euclidean distance between input vector and class vector
def compute_distance(input_vector, class_vector):
    return sqrt(
        pow(float(input_vector[0]) - float(class_vector["Sepal Length"]), 2) +
        pow(float(input_vector[1]) - float(class_vector[" Sepal Width"]), 2) +
        pow(float(input_vector[2]) - float(class_vector[" Petal Length"]), 2) +
        pow(float(input_vector[3]) - float(class_vector[" Petal Width"]), 2)
    )

# Define the classification function
def classify_iris(input_vector):
    # Access the broadcasted class vectors
    class_vectors = broadcast_class_vectors.value
    class_vector_1 = class_vectors[0]
    class_vector_2 = class_vectors[1]
```

✓ 3s completed at 12:27 AM

Part 2: SQL Notebook in Databricks

Do this exercise in a SQL notebook of databricks. Here let us do some OLAP exercises in Spark-SQL. Exercises here are inspired by an OLAP demo[1] from Jennifer Widom, Stanford University. It works on the same star schema discussed in the lectures. DDL and insert scripts are available here.

Perform the following operations on this dataset.

Q1. Create tables for the schema as defined in DDL script and add data using the INSERT scripts for all tables.

Code and Output:

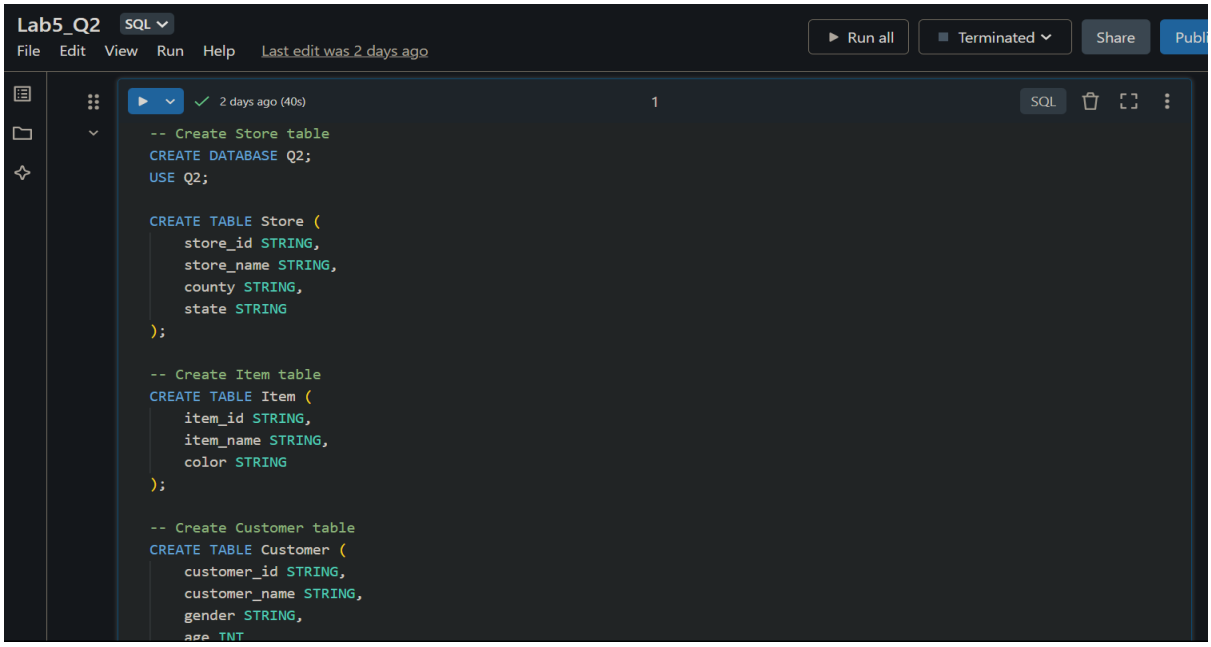

Q2: Create a Cube that Stores Sales Amount summaries. Let this Cube be materialized as a relational table and named as SalesCube.

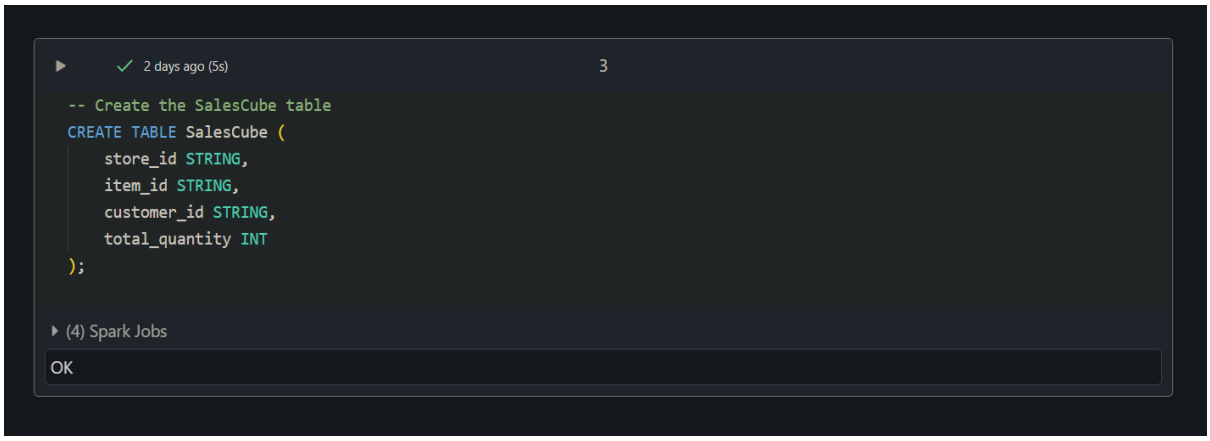

Q-3: List total sales (amount) of all items of category 'Tshirts' Project: item , sum(amount)

Code and Output:

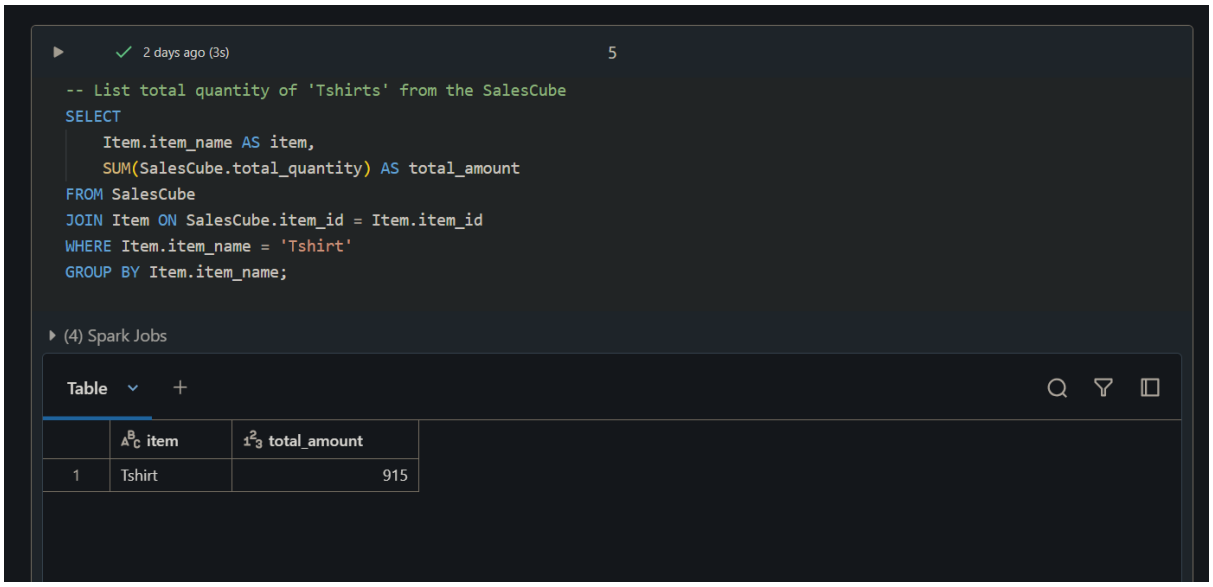

Q-4: List total sales (amount) of all items of category 'Tshirts' at 'store1' Project: item , sum(amount)

Code and Output:

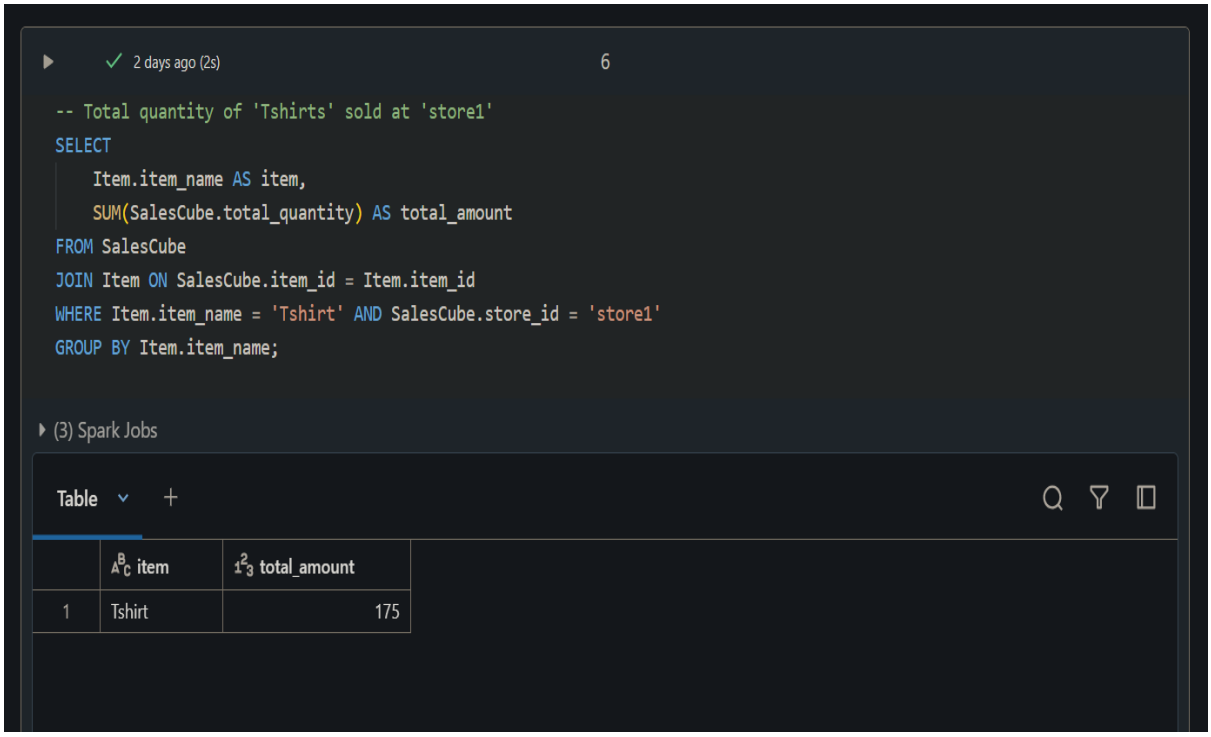

Q5: Give total sales (amount) of customer 'cust1' . Project: sum(amount)

Code and Output:

Q6: Give total sales (amount) of customer 'cust1' from items of category 'Tshirts' Project: item , sum(amount)

Code and Output:

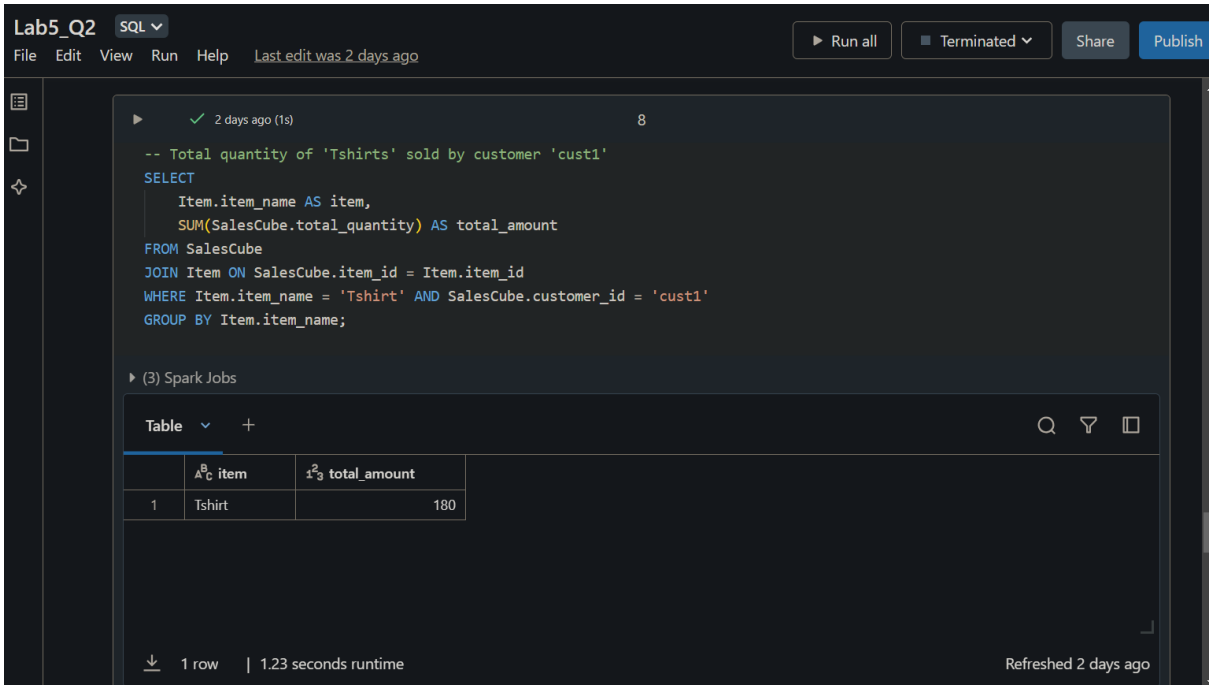

Q7: Give total sales (amount) of all stores of 'WA' state for items of category 'Tshirts' Project: store , item , sum(amount)

Code and output:

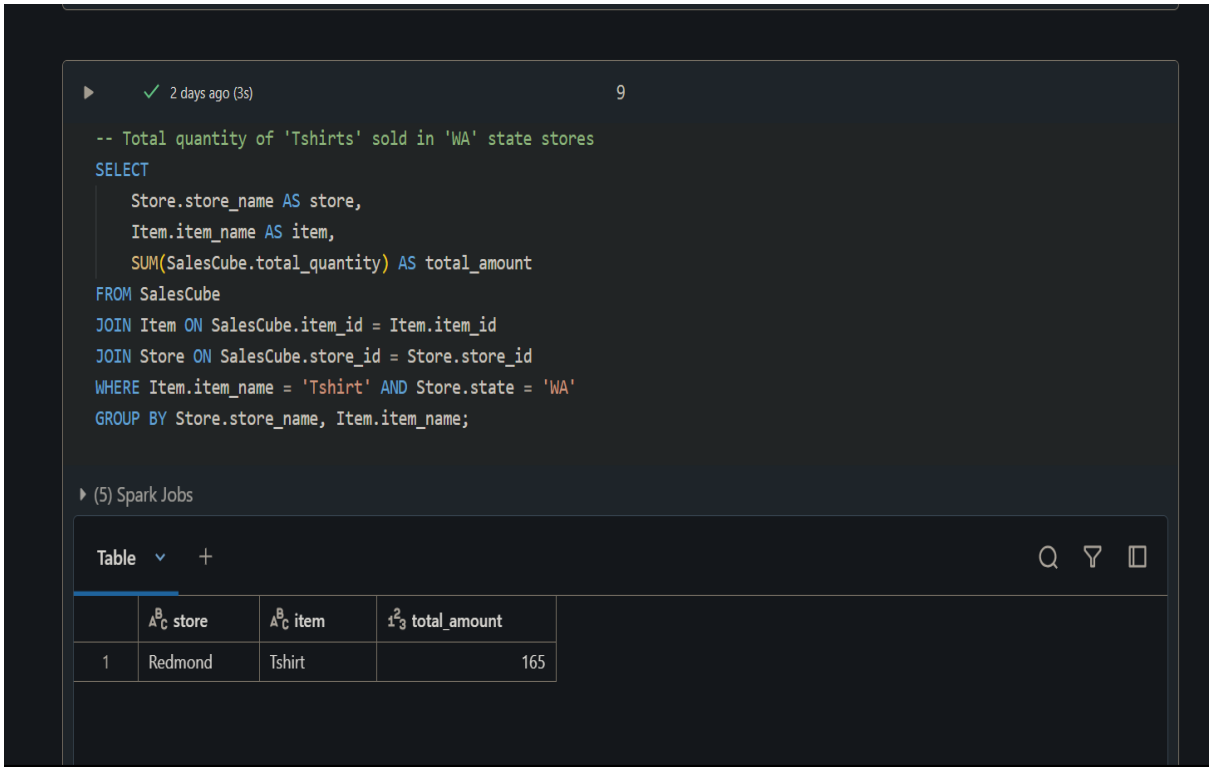

Q8: Give gender-wise total sales (amount) of items in the 'Jacket' category. Project: gender , item , sum(amount)